

# Coding and Configuration for Secure Systems

John O'Keefe-Odom



# Contents

|        |                                                                                                                                          |    |
|--------|------------------------------------------------------------------------------------------------------------------------------------------|----|
| 1      | Amnesiac System Planning, Experimentation, and Implementation in FreeBSD                                                                 | 1  |
| 1.1    | Introduction                                                                                                                             | 1  |
| 1.1.1  | Hardware Configuration                                                                                                                   | 1  |
| 1.1.2  | Intentions: Paving the Road to Jails                                                                                                     | 2  |
| 1.2    | Methodologies                                                                                                                            | 3  |
| 1.2.1  | Preparing the Baseboard                                                                                                                  | 3  |
| 1.2.2  | Incremental Advancement to Amnesiac System                                                                                               | 3  |
| 1.2.3  | Identifying and Allocating ZFS Drive Qualities                                                                                           | 3  |
| 1.2.4  | What are we going to want the ZFS kit to do? How are we going to use the drives? How can we wrestle with this new type of configuration? | 4  |
| 1.2.5  | What We Can Learn From Installing an OS                                                                                                  | 6  |
| 1.2.6  | Basic Operating System Installation Goals                                                                                                | 6  |
| 1.2.7  | Main Steps for FreeBSD Installation from <code>bsdinstall</code>                                                                         | 6  |
| 1.2.8  | No Root Password                                                                                                                         | 7  |
| 1.2.9  | Be Able to Build the Keys                                                                                                                | 7  |
| 1.2.10 | Partition, Mount, and Use a Thumbdrive from Live Disc's Shell                                                                            | 8  |
| 1.2.11 | Start the Box and the Live Disc's Shell                                                                                                  | 8  |
| 1.2.12 | Be Able to Plug in the USB Thumbdrive and Find Its Logical Address                                                                       | 8  |
| 1.2.13 | Be Able to Partition the USB                                                                                                             | 8  |
| 1.2.14 | Be Able to Write to the USB                                                                                                              | 9  |
| 1.2.15 | Be Able to Put Some Key Pairs on the USB                                                                                                 | 9  |
| 1.2.16 | Be Able to Use Commands to Build Out the Drives                                                                                          | 9  |
| 1.2.17 | Goals for Amnesiac Systems                                                                                                               | 11 |
| 1.3    | Results                                                                                                                                  | 11 |
| 1.3.1  | Configure for SWAP Disk, ZFS Cache, ZFS SLOG Mirror, and ZFS RAIDz2                                                                      | 12 |
| 1.3.2  | Startup with <code>boontononly.iso</code>                                                                                                | 12 |
| 1.3.3  | Mount the USB to Save the Metadata Backup File                                                                                           | 13 |
| 1.3.4  | Create Swap Disk in Bay 0                                                                                                                | 13 |
| 1.3.5  | Create ZFS Cache in Bay 1                                                                                                                | 13 |
| 1.3.6  | Create ZFS SLOG in Bays 2 and 3                                                                                                          | 13 |
| 1.3.7  | Create ZFS Work Disks for Jails                                                                                                          | 14 |
| 1.3.8  | Bay 4                                                                                                                                    | 14 |
| 1.3.9  | Bay 5                                                                                                                                    | 14 |
| 1.3.10 | Bay 6                                                                                                                                    | 14 |
| 1.3.11 | Bay 7                                                                                                                                    | 14 |
| 1.3.12 | Encrypt those partitions                                                                                                                 | 15 |

|          |                                                 |    |
|----------|-------------------------------------------------|----|
| 1.3.13   | Commands to Encrypt the Partitions . . . . .    | 15 |
| 1.3.14   | To Preserve the Metadata Backup Files . . . . . | 16 |
| 1.3.15   | File Name Check: Editing Labels . . . . .       | 17 |
| 1.3.16   | To Use the Swap Disk . . . . .                  | 17 |
| 1.3.17   | Create the zpool for BLONDIE . . . . .          | 19 |
| 1.3.18   | Create a mountpoint . . . . .                   | 19 |
| Endnotes |                                                 | 21 |
| 2        | FreeBSD Tape Operations                         | 23 |
| 2.1      | Introduction . . . . .                          | 23 |
| 2.2      | Methodology . . . . .                           | 23 |
| 2.3      | Results . . . . .                               | 23 |
| 2.3.1    | Code Notes . . . . .                            | 23 |
| 2.3.2    | Sample Files: test.txt . . . . .                | 24 |
| 2.3.3    | Sample Files: test2.txt . . . . .               | 24 |
| 2.3.4    | Sample Files: test3.txt . . . . .               | 24 |
| 2.3.5    | Sample Files: test4.txt . . . . .               | 24 |
| 2.3.6    | Sample Files: test5.txt . . . . .               | 24 |
| 2.3.7    | Sample Script: scriptedDemo.sh . . . . .        | 24 |
| 2.4      | Analysis . . . . .                              | 27 |
| 2.4.1    | Quality . . . . .                               | 27 |
| 2.4.2    | Tests . . . . .                                 | 27 |
| 2.5      | Discussion . . . . .                            | 27 |
| 2.6      | References . . . . .                            | 28 |
| Endnotes |                                                 | 29 |

# 1 Amnesiac System Planning, Experimentation, and Implementation in FreeBSD

In this technical note, we explore configurations of the FreeBSD operating system. With a goal of creating an amensiac system that keeps the operating system separate from working storage drives, we experiment with various configurations. We demonstrate the planning of drive configuration coding and the reasoning behind the choices. First, we code a system that uses the operating system on the storage drive discs. Then we evolve the plan into using a live cd for operating system use with drives that are connected to the immutable system in RAM via configuration code. Using an amensiac system like this, we would be able to reset the operating system to a known immutable state with every startup. The drives could be disconnected and reattached to other systems for various purposes. Throughout, we cover specific commands for drive and operating system management and operation.

## 1.1 Introduction

We're starting with a Dell T610 Server. Made in 2011, this machine is about ten years old. We purchased this at a government e-waste auction. She needed some cleanup. We started on the front lawn, blowing out the dust with a leafblower.

Most of the time servers from these auctions will have minimal RAM. Usually the drives and any good accessories will be stripped out. Anytime we get a hard drive, from any source, one of the first actions we take is to erase the drives to NIST 800-88 Purge with dedicated devices. That kind of erasing can be done with dd in Kali or FreeBSD; but, the drive erasers usually have a verification feature that's helpful.

### 1.1.1 Hardware Configuration

This is my favorite computer. I've named her "Blondie" after Debbie Harry's New Wave band.

Now, she's configured with:

- A DVD RW drive
- A LTO III tape backup unit
- In her 8 drive bay, we have:
  - Two Velociraptor 80GB 10000 rpm drives.
    - ★ We'll use these in a mirror configuration to hold the main OS.
  - Two refurb 120GB SSDs
    - ★ We'll use this in a mirror as ZFS log drives.
    - ★ They'll help to make writes faster.
  - Four 500GB hard drives
    - ★ We'll use this as a RAIDZ-2 array to hold the jails.
    - ★ We have some smaller SAS drives.
    - ★ We have some SATA drives.

- We replaced the one CPU with two Xeon 5550 quad cores and added a cooling tower.
- We added more RAM. Currently, we're at 196GB.
- We added some more NICs.
  - We have four for the jails.
  - We have two for the OS.
  - We'll leave the two on the motherboard mostly unused, except for our iDRAC connection.
    - ★ iDRAC is a proprietary baseboard controller.
    - ★ We could use it to power cycle the machine.
    - ★ There's a web interface that let's us know about motherboard-related facts.
    - ★ We cannot give a GELI password through the console in our current setup.
- The server has two 500W power supplies.

Because the dual power supplies draw about 500 Watts apiece, this server gets its own circuit.

The drives we're using are older, but we got enough copies of the drives to be able to support backups. When making physical backups of the drives, it can be helpful to have other volumes of the same model. So, instead of getting the best available, we planned for the best we could afford at quantity.

The server does not have:

- Any type of graphics card or GPU
- Sound cards or microphone hardware
- Wireless networking or Bluetooth.

### 1.1.2 Intentions: Paving the Road to Jails

We'll use the box for development. It'll be our main development server. We'll use it with programs that we compile on a nearby poudriere package building box. In those package builds, we have segregated our builds into a pattern of jails that mimic the types of jails that will be needed on most of our computers.

So, there is a jail for functional groups like: development tools, offensive and defensive security programs, browsers, desktop software, and so on. Poudriere jails can be filled with whatever sets of programs we like. For successful builds, we find it's best to keep the jails small.

There will be three main groups of drives:

- an OS mirror pair
- a ZFS log mirror pair
- a ZFS RAIDZ-2 kit of four drives.

The OS pair will hold ideas like:

- jail configuration
- main host networking and VNets
- host-based user accounts and configuration.

The ZFS kits will hold disconnected jails. We'll set up a system of hierarchical

jails that provide us with a development environment made up of modular jails that will be the children of an overarching parent. These jails will communicate with each other and the parent through some custom VNETting. We'll do all of our jail scripting, controls, and configuration manually. Likewise, the base ZFS setup will be done with simple scripts.

## 1.2 Methodologies

### 1.2.1 Preparing the Baseboard

The T610 has a built-in RAID card that can't be totally bypassed. This means that since we'll want ZFS to control the drives as much as possible, we'll need to tell the hardware RAID to buzz off. We won't want hardware RAID to participate in the process. We'll just use the software RAID through ZFS.

We can't just unplug the RAID card. We tried that, and it leads to malfunctions on this machine. So, as a compromise, what we'll tell the machine is that each drive is its own VDEV and it's RAID-0. We'll end up with as many VDEVs as we have physical drives.

To prepare for a later step in this process, we'll note the serial number of each drive we load into position, its model number, and its capacity. We'll need to know that information, and we won't want to be hassled by having to power cycle the machine in the middle of the setup process just to check on those details.

Later on, we'll set VDEVs with ZFS; but, from this point on, we'll basically ignore the RAID card's description of these devices as VDEVs.

With the baseboard prepared, and the BIOS set to allow us to boot from disc, we'll get started with a plain install of FreeBSD to the first two drives in a mirror configuration.

### 1.2.2 Incremental Advancement to Amnesiac System

With all of the manual coding and configuration needed to set up an amnesiac system like the one we desired, it was a big change from the normal BSD installs we had done for several years. Our first pass at building an amnesiac system decayed into manually scripting a regular operating system installation that had most of the features we desired in the amnesiac system. Both the on-disc OS system and the amnesiac system are detailed below.

### 1.2.3 Identifying and Allocating ZFS Drive Qualities

We need to be able to detect the drives and partitions that we want to work with using `geom`, `gpart`, and `zfs` commands. If, for some reason, you can't find those drives, then you need to address that. While working with this particular computer, I've found that the leading cause of failing to detect a drive goes back to the VDEV assignment in the hardware RAID. If the drive were pulled, this particular RAID card

Table 1.1: ZFS Concepts and Commands

| Topic             | Concept                                                      | Command phrase example                          |
|-------------------|--------------------------------------------------------------|-------------------------------------------------|
| Disk labels       | Use the label to associate the logical with the physical     | [Host] - [PositionNumber 0-7] - [Serial_Number] |
| ZFS Compression   | Faster performance: less space                               | zfs set compress=gzip-6 [DIRECTORY]             |
| ZFS zpool comment | Make a note on the pool                                      | zfs set comment="- "                            |
| quota             | Limit growth                                                 | -                                               |
| reservation       | Set aside space for directory and its children               | -                                               |
| refreservation    | Set aside space but don't count snapshots and child datasets | -                                               |
| canmount          | Control mounting                                             | zfs set canmount=off [DIRECTORY]                |
| mountpoint        | Specify where to mount                                       | zfs set mountpoint=/SOME [DIRECTORY]            |
| readonly          | Control readonly                                             | zfs set readonly=off [DIRECTORY]                |
| exec              | Control if it can execute                                    | zfs set exec=off [DIRECTORY]                    |
| copies            | Make more than one copy                                      | zfs set copies=2 [DIRECTORY]                    |

will need to have those configuration changes addressed. Often, this means rebooting the machine. It could be done through the iDRAC baseboard controller in some models; but the one we have here is not configured to do those tasks.

1.2.4 What are we going to want the ZFS kit to do? How are we going to use the drives? How can we wrestle with this new type of configuration?

After reading over FreeBSD Mastery: ZFS by Alan Jude and Michael Lucas, we came to a few conclusions. We chose to use RAIDZ-2 with Log drives because of the number of drives we have. Some of the topics covered in that book that we'll want to apply include ideas in Table 1.1. In our first iterations we decided to install the OS on the hard drives in a pattern that looks like a conventional installation. Once the drive partitioning and utilization commands were worked out, we then started over and undertook the amnesiac version with the OS on a live disc that runs with attached drives.



Table 1.2: ZFS Main OS Use Hard Disk: Initial Estimates

|                                    | Estimate for 80 GB HDD |
|------------------------------------|------------------------|
| Blank: 15% "Free Space"            | 12 GB                  |
| Swap: Under-allocate, deliberately | 10 GB                  |
| ZFS Main Partition: Remainder      | 58 GB                  |

Table 1.3: ZFS Jail Use Hard Disk: Initial Estimates

|                                     | Estimate for 500 GB HDD |
|-------------------------------------|-------------------------|
| Blank: 15% "Free Space"             | 75 GB                   |
| Swap: 25% Physical RAM (196 GB max) | 50 GB                   |
| ZFS Main Partition: Remainder       | 375 GB                  |

Our goal for the system will be to split up the drives into three main groups, as described in Tables 1.2, 1.3, and 1.4. On each drive, we'll also allow:

- 15% of the disk's own capacity as "Free Space" by allocating a blank partition
- Swap by 100 | 50 | 33 | 25 percent of RAM (with a 25% minimum for all disks) or similar
- Remainder as one large ZFS main partition.

We decided we would want those qualities in an amnesiac system, with the exception of running the OS from a live disc.

For the OS-On-Disk system, this should bring us:

- Enough room for the main OS
- A 4 disk RAIDZ-2 array of almost 1.49 TB of actual storage
- Fast swap and copy-on-write
- With each disk capable of expansion or change.

The unused space on the drives (a blank partition of 15%) may look like we're grossly under-utilizing the disks. We're planning for future emergencies. Do we need more swap? Maybe we can add it in. Do we need to expand or resize the disk? Maybe we can delete that blank "Free Space" partition and use that. By not trying to use everything, we're giving ourselves some room to solve future problems.

Table 1.4: ZFS SLOG Log Use SSD: Initial Estimates

|                                     | Estimate for 120 GB SSD |
|-------------------------------------|-------------------------|
| Blank: 15% "Free Space"             | 18 GB                   |
| Swap: 25% Physical RAM (196 GB max) | 50 GB                   |
| ZFS Log Partition: Remainder        | 52 GB                   |

## 1.2.5 What We Can Learn From Installing an OS

Later on, we'll install the other drives with `gpart` and `zfs` commands. This first pair, through, we could do with `bsdinstall`. If we choose that, though, then we must commit the whole disk to the process. That would be how we would normally stand up a disk drive. In this case, we'll conduct a manual process that mimics the main steps of `bsdinstall` to show that the installation can be done manually. Part of these same kinds of steps are critical tasks when administering and installing jails.

## 1.2.6 Basic Operating System Installation Goals

We're going to do an early phase, basic installation of the OS. At this point, we'll really only set up a pair of users and some plain vanilla settings for the OS. We want to be able to:

- to establish our ZFS kit
- to configure our main jail relationships
- to be able to check our networking with other hosts
- to set the basic operating system to call packages from our package server.

We can come back later and make a more sophisticated environment. At this stage, we will just lay the foundation for that work. If we don't get the ZFS kit stood up, then our more complex work with jails won't have an environment in which to thrive.

## 1.2.7 Main Steps for FreeBSD Installation from **bsdinstall**

We can see that `bsdinstall` is broken up into about a dozen main segments. They are:

- Keyboard mapping
- Hostname
- Networking configuration
- Wireless networking
- Partitioning
- Mounting temporary root
- Fetch the OS
- Extract the OS
- Set a root password
- Set the time
- Add users
- Enable some services
- Copy in the config [2].

We're not going to do all of these. We'll start with partitioning. We won't do any wireless networking or set a root password. We're going to skip the wireless networking for now because the hardware doesn't have that capability. In the newer versions of `bsdinstall`, there are some protective measures not listed above. We won't tackle those just yet.

## 1.2.8 No Root Password

We won't set a root password because we will expect to always ssh into the root account with a key. By not setting a root password, we will create a situation where a key is always required. This can be dangerous in that it can create a single point of failure for the system.

Another risk is that we will require the ability to ssh in as root. Later, as we establish software firewalls on the host, we will want to limit what addresses can be used to access this root with ssh. It's a little dangerous, but it's also a little different.

We have some risk controls that will be in place. First, we often store our keys in a special collection of USB drives. Next, we will have 2FA on the ssh that will require a TOTP (Time-based One Time Password). Finally, we will limit what's done in this drive. Since most of our programs and work will be done in the jails, there is a limited amount of influence that this part of the system will have. It's there to start the box and spin up the jails. If we ever have to totally destroy this installation and hook up our jails to another system, that is possible.

What do we gain from no root password? An extra measure of security. If there is no password, then password enumeration attempts should all fail. This not only goes for root, but it goes for other users. Many of the automated login attacks we've seen have been simple dictionary attacks; these brute force exercises rarely attempt to account for the types of login procedures that we expect to put in place.

Not having a root password may also truncate or stop our ability to use commands like `su` or `sudo`. As we're trying to stop privilege escalation, we'll also be stopping legitimate means of `priv-esc`. It could be inconvenient, but we've used this technique before.

Also, there is another risk: leaving the needed key lying around. For example, if we give the key to a user account operator, like ourselves, and then we import copies of the key to the box for easy use: well, then those keys could be used and re-used locally. This would be the equivalent for storing a plain-text password on the box. They keys are ultimately simple text files with decorated, long, strings of hashes. It's the software that goes with the keys, combined with the configuration of the programs for logging in to a service, that makes them special. Switching from passwords to keys is not a perfect answer; but, it's the answer we'll use this time.

If someone wishes to put a root password in place, then that's a procedure that will be easy to implement. For now, though, there will be no root password. And that means we have to build the keys we will use to log in.

## 1.2.9 Be Able to Build the Keys

We can build our keys using the live disc's shell. and a thumbdrive. The main command to build the keys is:

```
ssh-keygen -b 1024 -t rsa
```

Copy the files output to a thumbdrive and keep them handy for mounting.

### 1.2.10 Partition, Mount, and Use a Thumbdrive from Live Disc's Shell

Let's take some time to go over some detailed instructions on how to build those keys and get them onto a thumbdrive. The live disc's shell mostly uses a read-only file system. So, some people may not be familiar with how we are going to get the keys made and saved anywhere that will persist.

#### 1.2.11 Start the Box and the Live Disk's Shell

With the hardware booted to the LiveCD (installation disc), we select "Shell" from the three main choices in the first menu. This drops us into the shell. If we try to save a file of any kind, we'll probably get an error message back that says we're in a read-only file system. Our main set of directions for this is in the FreeBSD Handbook. We look carefully at the sections for USB storage and "Adding Disks"

#### 1.2.12 Be Able to Plug in the USB Thumbdrive and Find Its Logical Address

```
camcontrol devlist
```

This should show the USB drive. If the list is long, we can pipe it to more and slowly advance through the list. In our case, the USB was at da0.

The listing should also show any other storage device that's plugged in to the machine. So, all of those hard drives loaded into the bays should appear. If they don't, then take it as a warning sign that we're having a problem with the RAID card's physical drive VDEVs. On the T610, we were seeing drives in positions 0 to 7 in addition to the USB thumbdrive we're working on now.

#### 1.2.13 Be Able to Partition the USB

Since our USB has not been prepared for use with anything, we have to partition it. We can pretty much use the default commands from 17.2 "Adding Disks" to set up the drive. Here, we create a simple UFS file system and mount it to a directory that we'll put in /tmp. If our USB is already partitioned, we would want to skip this part.

```
# gpart create -s GPT da0
# gpart add -t freebsd-ufs -a 1M da0
# newfs -U /dev/da0p1
```

Mount the USB to the /tmp Directory:

```
# mkdir /tmp/someDir
# mount /dev/da0p1 /tmp/someDir
```

### 1.2.14 Be Able to Write to the USB

We could now write a simple text file onto /tmp/someDir, unmount the drive, and then re-mount it to check our work. That would go like:

```
# cd /tmp
# umount /tmp/someDir
# rm -r /tmp/someDir
```

Unplug the drive. Plug it in another slot. Use camcontrol devlist to find it again. It'll probably get the same address, like "da0".

```
# mkdir /tmp/otherDir
# mount /dev/da0p1 /tmp/otherDir
# ls /tmp/otherDir
```

### 1.2.15 Be Able to Put Some Key Pairs on the USB

Using these same principles, we can mount our USB thumbdrive, generate ssh keys, and save them there for later.

```
# mkdir /tmp/otherDir/keys
# ssh-keygen -b 1024 -t rsa
```

During the ssh-keygen dialog, simply direct the program to save the keys to /tmp/otherDir/keys. Be sure to record your passphrase somewhere.

### 1.2.16 Be Able to Use Commands to Build Out the Drives

We'd like to show a quick overview of some of the drive-related commands. These will be applied soon.

To list what devices are detected by the OS, we can use:

```
gpart show
```

We will want to note the serial number of each drive, by position. We then note which drives we want to use for the main body of the RAIDZ-2 kit. All of this will become one pool of ZFS VDEVs. Then we'll add in the log mirrors. So, in the sequence of 0-7, numbering the drives, we should expect to skip those two for now.

Make the notes if you haven't already because we're going to list all of the drives for the main pool in one pop. This task may have been done before we set up the physical VDEVs for the hardware RAID.

We're going to choose to put a swap partition on the SSDs. We can add one onto the other disks, but a maximum of 4 swap partitions is recommended in the FreeBSD Handbook. Later on, we will be able to select swap partitions by mounting

them and using commands like `swapon` `swapoff`.

To put the ZFS partition on the remaining disks:

```
gpart create -s SCHEME DISK_DEVICE
```

```
gpart create -s gpt da4
```

To add the main ZFS partition:

```
gpart add -a 1m -l ZFS_NUMBER -t freebsd-zfs DISK_DEVICE
```

```
gpart add -a 1m -l ZFS-BLONDIE-4-WD123465789 -t freebsd-zfs  
da4
```

To add the partition's swap:

```
gpart add -a 1m -slg -l SWAP_NUMBER -t freebsd-swap DISK_DEVICE
```

```
gpart add -a 1m -slg -l SWAP-BLONDIE-4-WD123465789 -t  
freebsd-swap da4
```

To add a logging partition on the SLOG SSDs:

```
gpart add -a 1m -l ZLOG_NUMBER -t freebsd-zfs DISK_DEVICE
```

```
gpart add -a 1m -l ZLOG-BLONDIE-2-WD123465789 -t freebsd-zfs  
da2
```

With each disk partitioned, we can add them to the pool for our RAIDZ-2 kit.

```
zpool create RAID POOL_NAME DISK_LABEL/ZFS_PARTITION_LABEL  
...next disk/partition...
```

```
zpool create BLONDIE-MAIN raidz2 \  
BLONDIE-4-WD12345/ZFS-BLONDIE-4-WD12345 \  
BLONDIE-5-WD12345/ZFS-BLONDIE-5-WD12345 \  
BLONDIE-6-WD12345/ZFS-BLONDIE-6-WD12345 \  
BLONDIE-7-WD12345/ZFS-BLONDIE-7-WD12345
```

We can add on the mirrors for logging:

```
zpool add BLONDIE-MAIN log mirror \  
da2
```

BLONDIE-2-WD12345/ZLOG-BLONDIE-2-WD12345 \  
BLONDIE-3-WD12345/ZLOG-BLONDIE-3-WD12345

Once we get to the point that we have an operating system on disks, and disks ready for use with logs and jails, then we're at the end of this phase of system administration.

### 1.2.17 Goals for Amnesiac Systems

We chose at this point to do some experimental installs. For clarity, notes on those are not presented here. We did change our goals for the system. We recognized that we wanted:

- Use a LiveCD for an immutable OS
- Encrypt disc partitions instead of the whole disc
- Ensure we save keys and metadata to promote disc recovery
- Prepare commands in advance to reduce complexity.

From our testing, we recognize that a Live CD offered the best answer for an immutable system. It does come at the price of some instability. The "El Torrito" discs have a limited set of services. Adjusting the configuration of the operating system means drafting, burning, and testing a new disc. The "shell" instance on a FreeBSD installation disc is reportedly less stable; but, we did not know why or how it was less stable than a normal operating system. We worked through our disc commands and kept careful notes. Adding the physical serial number of the disc alone increased the complexity of manually typing in each command. We discovered the need to encrypt the discs by partitions instead of the whole disc because there was a visibility problem during the execution of discovering and attaching the discs. There was a poorly documented detail about protecting the disc metadata and keys. The keys are provided during the setup dialog; but, these highly important items will disappear after the setup process closes. This means it is important that an operator recognize the need to intercept and to record these keys. Without them, a later recovery of a disc is impossible. Since one typo would derail our experimental setups, we had to keep a second support computer on hand to record notes and consult references.

## 1.3 Results

As our experimentation concluded, we had learned enough to stand a reasonable chance to setup an amnesiac machine. Its features include:

- An immutable OS
- Encrypted disc partitions
- Swap and logging for ZFS disc features
- Metadata and keys saved to a separate USB drive
- A disc for holding FreeBSD jails (containerization)
- Enough room to expand, to update, and to recover.

Table 1.5: FreeBSD Swap Disk – Estimate of 500 GB HDD

|                          |        |
|--------------------------|--------|
| GEOM Disk List Mediasize | 465 GB |
| Blank: 15% "Free Space"  | 70 GB  |
| Swap                     | 395 GB |

Table 1.6: ZFS Cache (Read Cache) Disk – Estimate of 80 GB HDD

|                          |       |
|--------------------------|-------|
| GEOM Disk List Mediasize | 74 GB |
| Blank: 15% "Free Space"  | 11 GB |
| Swap                     | 63 GB |

1.3.1   Configure for SWAP Disk, ZFS Cache, ZFS SLOG Mirror, and ZFS RAIDz2

Changed config to keep the RAIDZ2 for main work. I was having trouble with the SSDs; so, I removed them and replaced them with two 80GB Western Digital Velociraptors. I added another WD Velociraptor for cache. Then we added a Western Digital Blue 500GB for a swap disk. That sounds like a lot of swap, but we hope to upgrade to near 196GB of RAM this year. So, a little more than twice that would be 400GB of swap. The next disk size on hand was 500GB.

We had some mis-aligned partitions with RAIDZ2 work for Blondie. We were using the full 500GB on the label as our actual amount of media to work with. Since less is available, we will have to set those up again.

We used gpart delete to delete each partition. Then we used gpart destroy to delete the partition table for the whole disk. In the end, gpart showed us only the CD.

1.3.2   Startup with **boontonly.iso**

For this run, we'll begin by using the boot-only disk. We won't plan to ever boot from these hard drives. We'll just use the shell on the CD. As with the others, we'll mount a USB to store metadata information. We'll partition and encrypt the disks. In the following examples, we provide the actual serial numbers for the hard drives

Table 1.7: ZFS SLOG Log Mirror (Write Cache) Disk – Initial Estimates

|                                     |                         |
|-------------------------------------|-------------------------|
|                                     | Estimate for 120 GB SSD |
| GEOM Disk List Mediasize            | 465 GB                  |
| Blank: 15% "Free Space"             | 18 GB                   |
| Swap: 25% Physical RAM (196 GB max) | 50 GB                   |
| ZFS Log Partition: Remainder        | 52 GB                   |



Table 1.8: ZFS Jail Use Hard Disk – Estimate of 500 GB HDD

|                                     | Estimate for 500 GB HDD |
|-------------------------------------|-------------------------|
| GEOM Disk List Mediasize            | 465 GB                  |
| Blank: 15% "Free Space"             | 75 GB                   |
| Swap: 25% Physical RAM (196 GB max) | 50 GB                   |
| ZFS Main Partition: Remainder       | 375 GB                  |

that were used. This example was derived from notes we maintained during experimentation; serial numbers in other instances will vary, of course.

1.3.3 Mount the USB to Save the Metadata Backup File

```
mkdir /tmp/usb3
mount /dev/da0p1 /tmp/usb3
```

1.3.4 Create Swap Disk in Bay 0

```
gpart create -s GPT mfid0

gpart add -a 5K -s 395G -t freebsd-swap -l ZFS-BLONDIE-0-SWAP
mfid0
```

1.3.5 Create ZFS Cache in Bay 1

```
gpart create -s GPT mfid1

gpart add -a 5K -s 63G -t freebsd-zfs -l ZFS-BLONDIE-1-CACHE
mfid1
```

1.3.6 Create ZFS SLOG in Bays 2 and 3

```
gpart create -s GPT mfid2

gpart add -a 5K -s 63G -t freebsd-zfs -l ZFS-BLONDIE-2-SLOG
mfid2

gpart create -s GPT mfid3

gpart add -a 5K -s 63G -t freebsd-zfs -l ZFS-BLONDIE-3-SLOG
mfid3
```

### 1.3.7 Create ZFS Work Disks for Jails

#### 1.3.8 Bay 4

```
gpart create -s GPT mfid4
```

```
gpart add -a 5K -s 50G -t freebsd-swap -l ZFS-BLONDIE-4-SWAP-WCC2  
mfid4
```

```
gpart add -s 345G -t freebsd-zfs -l ZFS-BLONDIE-4-FILES-WCC2EE134  
mfid4
```

#### 1.3.9 Bay 5

```
gpart create -s GPT mfid5
```

```
gpart add -a 5K -s 50G -t freebsd-swap -l ZFS-BLONDIE-5-SWAP-WMC2  
mfid5
```

```
gpart add -s 345G -t freebsd-zfs -l ZFS-BLONDIE-5-FILES-WMC2E3914  
mfid5
```

#### 1.3.10 Bay 6

```
gpart create -s GPT mfid6
```

```
gpart add -a 5K -s 50G -t freebsd-swap -l ZFS-BLONDIE-6-SWAP-WCC2  
mfid6
```

```
gpart add -s 345G -t freebsd-zfs -l ZFS-BLONDIE-6-FILES-WCC2ETT39  
mfid6
```

#### 1.3.11 Bay 7

```
gpart create -s GPT mfid7
```

```
gpart add -a 5K -s 50G -t freebsd-swap -l ZFS-BLONDIE-7-SWAP-WCAY  
mfid7
```

```
gpart add -s 345G -t freebsd-zfs -l ZFS-BLONDIE-7-FILES-WCAYU2878  
mfid7
```

We should note that “Bay 7” is an observation of the physical hardware. The device (“mfid7”) could have whatever number assigned at the end that the machine gave it. If the hardware raid had a problem with recognizing a device, the number might not correlate to what bay it is in. When powering up, check your geometry.

### 1.3.12 Encrypt those partitions

During this phase, we'll encrypt the partitions. That means that we should have some recordkeeping materials on hand. Use a password manager or a notebook to carefully record the password that is about to be typed in. With each one completed, the machine will tell us that it has saved a file of backup metadata. We will want to copy that metadata file to the thumbdrive. If the encryption gets jammed up, or if there is some type of problem, then the metadata file will be used as part of the recovery process. Without a metadata file, recovering an encrypted partition can be difficult or impossible.

If you lose the metadata file, but still have routine access to the geli encryption because you know the password and the machine is functioning alright, then you can export another copy of the metadata. This means that anyone with root access to the machine, who can use geli commands, would also be able to export such a file. Keep this in mind if you need to secure your machine.

When we give the subcommand to "init," the machine will ask us to give the new password twice. This is part of the password setting process. When we give the subcommand to "attach," the machine will ask us to give the password as part of the normal attachment process. Type carefully, and look at the records carefully, because this is the main chance to make sure the password notes are right.

### 1.3.13 Commands to Encrypt the Partitions

```
geli init -l 256 mfid0p11  
geli attach mfid0p1
```

```
geli init -l 256 mfid1p1  
geli attach mfid1p1
```

```
geli init -l 256 mfid2p1  
geli attach mfid2p1
```

```
geli init -l 256 mfid3p1  
geli attach mfid3p1
```

```
geli init -l 256 mfid4p1  
geli attach mfid4p1
```

```
geli init -l 256 mfid4p2  
geli attach mfid4p2
```

```
geli init -l 256 mfid5p1  
geli attach mfid5p1
```

```
geli init -l 256 mfid5p2  
geli attach mfid5p2
```

```
geli init -l 256 mfid6p1
geli attach mfid6p1

geli init -l 256 mfid6p2
geli attach mfid6p2

geli init -l 256 mfid7p1
geli attach mfid7p1

geli init -l 256 mfid7p2
geli attach mfid7p2
```

#### 1.3.14 To Preserve the Metadata Backup Files

Let us reiterate: preserving these backups at the time of disk creation is essential for later recoveries. If these files are not saved, then a powerdown of the system will cause them to be lost. They're automatically stored in an ephemeral directory. Conserve the ability to recover a disk later by recording these metadata backup files to a safe location.

Each metadata backup file was probably created with a name that does not look like the drive's label. Instead, the machine assigns a backup file name based on geometry. Since we can have a hardware configuration change that could mess up that pattern, we will want to save our backup files with names that look more like the disk labels. That way, if the drives move, then we can better discover which metadata file goes with it.

```
cp /var/backups/mfid0p1.eli /tmp/usb3/backups
_27SEP2020/ZFS-BLONDIE-0-SWAP-WCC2ERZ56512.eli

cp /var/backups/mfid1p1.eli /tmp/usb3/backups
_27SEP2020/ZFS-BLONDIE-1-CACHE-WXL1E40C1352.eli

cp /var/backups/mfid2p1.eli /tmp/usb3/backups
_27SEP2020/ZFS-BLONDIE-2-SLOG-WXL1E4067503.eli

cp /var/backups/mfid3p1.eli /tmp/usb3/backups
_27SEP2020/ZFS-BLONDIE-3-SLOG-WXC0CA9V8431.eli

cp /var/backups/mfid4p1.eli /tmp/usb3/backups
_27SEP2020/ZFS-BLONDIE-4-SWAP-WCC2EE134983.eli

cp /var/backups/mfid4p2.eli /tmp/usb3/backups
```

```
_27SEP2020/ZFS-BLONDIE-4-FILES-WCC2EE134983.eli  
  
cp /var/backups/mfid5p1.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-5-SWAP-WMC2E3914420.eli  
  
cp /var/backups/mfid5p2.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-5-FILES-WMC2E3914420.eli  
  
cp /var/backups/mfid6p1.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-6-SWAP-WCC2ETT39577.eli  
  
cp /var/backups/mfid6p2.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-6-FILES-WCC2ETT39577.eli  
  
cp /var/backups/mfid7p1.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-7-SWAP-WCAYU2878968.eli  
  
cp /var/backups/mfid7p2.eli /tmp/usb3/backups  
_27SEP2020/ZFS-BLONDIE-7-FILES-WCAYU2878968.eli
```

### 1.3.15 File Name Check: Editing Labels

We got through all that, and found that we had named a file incorrectly. During experimentation, we had a typo; BAY 4 should be WCC2EE134983 but that's not what we saw when we checked our work. Not the end of the world. We can move the .eli metadata file, and we can use gpart modify to change the label.

```
gpart modify -i 1 -l  
ZFS-BLONDIE-4-SWAP-WCC2EE134983 mfid4  
  
gpart modify -i 2 -l  
ZFS-BLONDIE-4-FILES-WCC2EE134983 mfid4
```

gpart show and geli status should show our drives as we expect them to be.

At this point, all of the drives should be partitioned and decrypted and attached. We can build a zpool for our ZFS file system. We can attach that swap drive for the benefit of our operating system right now.

### 1.3.16 To Use the Swap Disk

REF: <https://www.cyberciti.biz/faq/create-a-freebsd-swap-file/>

```
swapon mfid0p1.eli swapinfo -k
```

64G turned out to be the maximum capacity. We will need to fix up the swap better. We can fit 6 x 64G (384G) of swap partitions in our planned space. So, let's modify and add.

```
gpart resize -i 1 -s 64G mfid0
```

We would need to add, to init and to attach five more swap partitions. We also have to restore mfid0 because there's a problem with reading the metadata.

```
gpart recover mfid0
gpart status mfid0
geli attach mfid0p1
```

```
gpart add -s 64G -t freebsd-swap -l
ZFS-BLONDIE-0-SWAP-B-WCC2ERZ56512 mfid0
```

```
gpart add -s 64G -t freebsd-swap -l
ZFS-BLONDIE-0-SWAP-C-WCC2ERZ56512 mfid0
```

```
gpart add -s 64G -t freebsd-swap -l
ZFS-BLONDIE-0-SWAP-D-WCC2ERZ56512 mfid0
```

```
gpart add -s 64G -t freebsd-swap -l
ZFS-BLONDIE-0-SWAP-E-WCC2ERZ56512 mfid0
```

```
gpart add -s 64G -t freebsd-swap -l
ZFS-BLONDIE-0-SWAP-F-WCC2ERZ56512 mfid0
```

Then we need to establish encryption for all of those.

```
geli init -l 256 mfid0p2
geli attach mfid0p2
geli init -l 256 mfid0p3
geli attach mfid0p3
geli init -l 256 mfid0p4
geli attach mfid0p4
geli init -l 256 mfid0p5
geli attach mfid0p5
geli init -l 256 mfid0p6
geli attach mfid0p6
```

using swapinfo, swapon and swapoff, we were able to remove some of the other swap partitions and add on these. Did see a kernel limit on swap. Total configured swap exceeded kern.maxswzone (98087976 pages). We could tune the

kernel, or we could reduce swap. For now, we reduce swap from 384GB to 324GB.

Similarly, we could add on the other swap partitions, although in our current use, we probably won't need them. They would serve us better in cases where we were doing more intensive work, or working on another machine. We can mount them now, anyway, to see that it can be done.

To add on all the other swap partitions:

```
swapon /dev/mfid4p1.eli  
swapon /dev/mfid5p1.eli  
swapon /dev/mfid6p1.eli  
swapon /dev/mfid7p1.eli  
swapinfo -h
```

### 1.3.17 Create the zpool for BLONDIE

We'll now create a mountpoint on the tmpfs, create the zpool, and then add in the cache and SLOG mirrored disks.

### 1.3.18 Create a mountpoint

Create a mountpoint so that we will be able to mount the zpool without errors as we create it.

```
mkdir /tmp/BLONDIE
```





## Endnotes

<sup>1</sup>"Geli(8)." Freebsd.org, 2025, [man.freebsd.org/cgi/man.cgi?query=geli&sektion=8&manpath=FreeBSD+10.3-RELEASE](http://man.freebsd.org/cgi/man.cgi?query=geli&sektion=8&manpath=FreeBSD+10.3-RELEASE). Accessed 10 June 2026.



## 2 FreeBSD Tape Operations

We explored the planning and use of LTO tape as a resource for storing backups. We describe common commands, patterns of MT and TAR commands. We recommend some techniques for testing, rehearsing, and using tape archives.

### 2.1 Introduction

We wanted to see if it would be possible for us to make use of an LTO tape machine that was on one of our salvaged computers. We'd often heard of the performance of using tape for backups. We'd seen some hints about the bravery needed for working with tape in modifying old mainframe programs. We rarely saw any strong tutorials or comprehensive directions for tape operations.

The average person has almost never used tapes as a part of personal computing these days. Businesses which do use tapes often do so as part of larger, commercial data center operations. For many of us, the closest we've come to tape operations was with using cassette tapes for storage on personal machines in the 1980s. LTO tapes are often praised but infrequently used. We wanted to break through this mystery.

### 2.2 Methodology

Our plan is to write a file to tape, eject the tape, and be able to read the file back from the tape again. Once there, we will want to do a simple sequence of file writes and reads without damaging the files on tape.

### 2.3 Results

We used a Dell T610 tower computer with an LT03 single deck tape machine on board. We have some LT03 tapes that can be used with it. Our tape drivers and common commands are built into FreeBSD.

Warning: We want to use the nsa driver. There is an sa driver, but it has some limitations. When we tried using sa, we were not able to properly advance down the tape through multiple file markers. Use nsa for the the driver.

#### 2.3.1 Code Notes

1. Place files in a directory that you control. Isolating the files to be copied to tape is safer than copying them directly from where they are.

2. `mt -f /dev/nsa0 load`

This will load the tape onto the machine.

3. `mt -f /dev/nas0 rewind`

This will spool the tape to the beginning.

4. `mt -f /dev/nsa0 status`

To check to see where we are on the tape. Assuming this is the first use of the tape, we could start writing immediately.

5. If there are files on the tape, we need to either scan ahead and log the files, or we need to consult our notes to find where on the tape we should go. We could easily overwrite a file and damage the sequence. Use the write persistent notch to protect the tape when possible.

6. `mt -f /dev/nsa0 com or eod`

To go to the end of the records.

7. `tar cvf /dev/nsa0 [FILE_TO_ARCHIVE]`

Once at the proper spot on the tape, give the command to archive the files.

8. `mt -f /dev/nsa0 weof1`

Write an end-of-file marker.

9. `tar -t`

To see the archive on the tape.

`tar xf /dev/nsa0`

To get the archive from the tape.

10. Write notes. Log and record what was put on the tape.

11. For the next record, just rewind and advance through the markers to get to the next blank space for writing a record.

### 2.3.2 Sample Files: test.txt

The quick brown fox jumped over the lazy dogs.

### 2.3.3 Sample Files: test2.txt

The other dog is still lazy, too.

### 2.3.4 Sample Files: test3.txt

The fox is just really fast.

### 2.3.5 Sample Files: test4.txt

Yet some more records.

### 2.3.6 Sample Files: test5.txt

And another record.

### 2.3.7 Sample Script: scriptedDemo.sh

```
#!/usr/bin/env sh
clear
echo "scriptedDemo"
```

```
echo "Press ENTER to continue."
read INPUT
### Check tape
echo "Check if tape is staged in machine."
echo "Press ENTER to continue."
read INPUT
clear
### Load the tape into the machine
echo "Here we go."
echo
echo "Loading tape into machine."
echo "mt -f /dev/nsa0 load"
mt -f /dev/nsa0 load
### Rewind tape
echo "Rewinding tape"
echo "mt -f /dev/nsa0 rewind"
mt -f /dev/nsa0 rewind
echo "Press ENTER to continue."
read INPUT
clear
echo "Reading status from tape."
echo "mt -f /dev/nsa0 status"
mt -f /dev/nsa0 status
echo
echo "Press ENTER to continue."
read INPUT
clear
echo "Looking for file 0 on the tape."
echo "We should be at the beginning."
echo "So, no mt commands this time."
echo "tar t -f /dev/nsa0"
tar t -f /dev/nsa0
echo
echo "Press ENTER to continue."
read INPUT
clear
echo "Looking for file 1 on tape."
echo "Rewind to the beginning."
mt -f /dev/nsa0 rewind
echo "Go forward space count 1 file."
echo "mt -f /dev/nsa0 fsf 1"
mt -f /dev/nsa0 fsf 1
echo "tar t -f /dev/nsa0"
tar t -f /dev/nsa0
echo "Press ENTER to continue."
```

```
read INPUT
clear
echo "Looking for file 2 on tape."
echo "Back up to 0 and go forward space count 2 files."
echo "mt -f /dev/nsa0 rewind"
mt -f /dev/nsa0 rewind
echo "mt -f /dev/nsa0 fsf 2"
mt -f /dev/nsa0 fsf 2
echo "Show where we are with status."
echo "mt -f /dev/nsa0 status"
mt -f /dev/nsa0 status
echo
echo "Look in that spot with tar."
echo "tar t -f /dev/nsa0"
tar t -f /dev/nsa0
echo "Press ENTER to continue."
read INPUT
clear
echo "Looking for third archive on tape."
echo "Back up to 0 and go forward space count 3 files."
echo "mt -f /dev/nsa0 rewind"
mt -f /dev/nsa0 rewind
echo "mt -f /dev/nsa0 fsf 3"
mt -f /dev/nsa0 fsf 3
echo "See where we are with status."
echo "mt -f /dev/nsa0 status"
mt -f /dev/nsa0 status
echo
echo "See what is here with tar."
echo "tar t -f /dev/nsa0"
tar t -f /dev/nsa0
echo "Press ENTER to continue."
read INPUT
clear
echo "Take the tape offline."
echo "Rewind tape."
echo "mt -f /dev/nsa0 rewind"
mt -f /dev/nsa0 rewind
echo "mt -f /dev/nsa0 offline"
mt -f /dev/nsa0 offline
exit 0
```

## 2.4 Analysis

### 2.4.1 Quality

1. We found there were significant weaknesses in tape sequencing. Unlike random access memory, overwriting can damage end of file markers which are only navigational delimiters readily usable with the tape. Simple practice showed it was very easy to accidentally overwrite a tape marker. All subsequent files on the tape might be inaccessible as a result.

2. Logging. Since there may be many files on these long tapes, logging support and file navigation are essential. The usual lapses in notation can make frequent use of the tapes more hazardous with each use.

3. Security. Unlike geli-encrypted hard drives, the main file protection scheme is through whatever encryption we apply to the files themselves. Tar has some compression algorithms, but encrypting files is another step altogether. Since the files may be in storage for a long time, having the backups encrypted on tape can be a critical step. This encryption would need to be applied before IV.B.7. tar or archiving the files to tape.

### 2.4.2 Tests

Writing to tape should be rehearsed. Checking status and reading files to confirm our location on the tape is a good idea.

Any write operation should be couched on a good, rehearsed read script. Isolate the files to be written to their own directory. Add an encryption step if necessary. Rehearse the tar tape archive commands. Compartmentalize the steps of the tape operations plan to reduce catastrophic failures and a loss of data.

## 2.5 Discussion

Tape operations have been long recommended as the preferred form of backup. Off-site backup offers strong advantages to an on-premises system operator. We can't count the number of times we've seen backups recommended; yet we find very few practical tutorials or directions.

Tape backups and hard drive backups might be easily and securely routed to an off site location through the US mail. They could be mailed to yourself at a personal mailbox, remote office, P.O. Box, or to a trusted person. Notice that any person or entity entrusted to possess the data may also be at risk or some sort of liability.

## 2.6 References

A. Books – none. 2023 versions of the FreeBSD handbook don't cover tape operations.

B. man pages – `mt`, `tar` man pages.

C. URLs – FreeBSD man pages for above. `man.freebsd.org/cgi/man.cgi?query=mt` `tar`

Text with<sup>1</sup> inline references. More text<sup>2</sup>



## Endnotes

<sup>1</sup>First note.

<sup>2</sup>Second note.